

Machine Learning Validation via Rational Dataset Sampling with **astartes**

Jackson W. Burns Kevin A. Spiekermann Himaghna Bhattacharjee
Dionisios G. Vlachos William H. Green

3 April 2023

Summary

Machine Learning (ML) has become an increasingly popular tool to accelerate traditional workflows. Critical to the use of ML is the process of splitting datasets into training, validation, and testing subsets that are used to develop and evaluate models. Common practice in the literature is to assign these subsets randomly. Although this approach is fast and efficient, it only measures a model’s capacity to interpolate. Testing errors from random splits may be overly optimistic if given new data that is dissimilar to the scope of the training set; thus, there is a growing need to easily measure performance for extrapolation tasks. To address this issue, we report **astartes**, an open-source Python package that implements many existing similarity- and distance-based algorithms to partition data into more challenging splits that can better assess out-of-sample performance. This publication focuses on use-cases within cheminformatics. However, **astartes** operates on arbitrary vector inputs, so its principals and workflow are generalizable to other ML domains as well. **astartes** is available via the Python package manager **pip** and is publicly hosted on GitHub (github.com/JacksonBurns/astartes).

Statement of Need

Machine learning has sparked an explosion of progress in chemical kinetics (Komp, Janulaitis, and Valleau 2022; Spiekermann, Pattanaik, and Green 2022a), drug discovery (X. Yang et al. 2019; Bannigan et al. 2021), materials science (Wei et al. 2019), and energy storage (Jha et al. 2023) as researchers use data-driven methods to accelerate steps in traditional workflows within some acceptable error tolerance. To facilitate adoption of these models, researchers must critically think about several topics, such as comparing model performance to relevant baselines, operating on user-friendly inputs, and reporting performance on both interpolative and extrapolative tasks. **astartes** aims to make it straightforward for machine learning scientists and researchers to focus on two important points: rigorous hyperparameter optimization and accurate performance evaluation.

First, **astartes**’ key function `train_val_test_split` returns splits for training, validation, and testing sets using an **sklearn**-like interface. This partitioning is crucial since best practices in data science dictate that, in order to minimize the risk of hyperparameter overfitting, one must only optimize hyperparameters with a validation set and use a held-out test set to accurately measure performance on unseen data (Ramsundar et al. 2019; Géron 2019; Lakshmanan, Robinson, and Munn 2020; Huyen 2022; Wang et al. 2020). Unfortunately, many published papers only mention training and testing sets but do not mention validation sets, implying that they optimize the hyperparameters to the test set, which would be blatant data leakage that leads to overly optimistic results. For researchers interested in quickly obtaining preliminary results without using a validation set to optimize hyperparameters, **astartes** also implements an **sklearn**-compatible `train_test_split` function.

Second, it is crucial to evaluate model performance in both interpolation and extrapolation settings so future users are informed of any potential limitations. Although random splits are frequently used in the cheminformatics literature, this simply measures interpolation performance. However, given the vastness of chemical space (Ruddigkeit et al. 2012) and its often unsmooth nature (e.g. activity cliffs), it seems unlikely that users will want to be restricted to exclusively operate in an interpolation regime. Thus, to encourage adoption of these models, it is crucial to measure performance on more challenging splits as well. The general workflow is: 1. Convert each molecule into a vector representation. 2. Cluster the molecules based on similarity. 3. Train the model on some clusters and then evaluate performance on unseen clusters that should be dissimilar to the clusters used for training. Although measuring performance on chemically dissimilar compounds/clusters is not a new concept (Meredig et al. 2018; Durdy et al. 2022; Jorner et al. 2021; Stuyver and Coley 2022; Tricarico et al. 2022; Terrones et al. 2023; Heinen, Rudorff, and Lilienfeld 2021; Bilodeau et al. 2023), there are a myriad of choices for the first two steps; our software incorporates many popular representations and similarity metrics to give users freedom to easily explore which combination is suitable for their needs.

Example Use-Case in Cheminformatics

To demonstrate the difference in performance between interpolation and extrapolation, we apply **astartes** to two relevant cheminformatics datasets and their corresponding tasks. For each dataset, a typical interpolative split is generated using random sampling. We also create two extrapolative splits for comparison. The first uses the cheminformatics-specific Bemis-Murcko scaffold (Bemis and Murcko 1996) as calculated by RDKit (Landrum and others, n.d.). The second uses the more general-purpose K-means clustering based on the Euclidean distance of Morgan (ECFP4) fingerprints using 2048 bit hashing and radius of 2 (Morgan 1965; Rogers and Hahn 2010). The QM9 dataset and RDB7 datasets were organized into 100 and 20 clusters, respectively. For each split, we create 5 different folds (by changing the random seed) and report the mean absolute error (MAE) and root-mean-squared error (RMSE). The values in Table 1 and Table 2 correspond to the mean \pm one standard deviation calculated across folds.

First is property prediction with QM9 (Ramakrishnan et al. 2014), a dataset containing approximately 133,000 small organic molecules, each containing 12 relevant chemical properties calculated at B3LYP/6-31G(2df,p). We train a multi-task model to predict all properties, with the arithmetic mean of all predictions tabulated below. Second is a single-task model to predict a reaction’s barrier height using the RDB7 dataset (Spiekermann, Pattanaik, and Green 2022b, 2022c). This reaction database contains a diverse set of 12,000 organic reactions calculated at CCSD(T)-F12 that is relevant to the field of chemical kinetics. Models were generated using a modified version of Chemprop (K. Yang et al. 2019) to train a deep message passing neural network to predict the regression targets of interest. We use the hyperparameters reported by Spiekermann, Pattanaik, and Green (2022a) as implemented in the **barrier_prediction** branch, which is publicly available on GitHub (Spiekermann et al., n.d.).

Table 1 and Table 2 show the expected trend in which the average testing errors are higher for the extrapolation tasks than they are for the interpolation task. The results from random splitting are informative if the model will be primarily used in interpolation settings. However, these errors are likely unrealistically low if the model is intended to make predictions on new molecules that are chemically dissimilar to those in the training set. Performance is worse on the extrapolative data splits, which present a more challenging task, but these errors should be more representative of evaluating a new sample that is out-of-scope. Together, these tables demonstrate the utility of **astartes** in allowing users to better understand the likely performance of their model in different settings.

Table 1: Average testing errors for predicting the 12 regression targets from QM9 (Ramakrishnan et al. 2014).

Split	MAE	RMSE
Random	2.02 ± 0.06	3.63 ± 0.21
Scaffold	2.20 ± 0.27	3.46 ± 0.49
K-means	2.48 ± 0.33	4.47 ± 0.81

Table 2: Testing errors in kcal/mol for predicting a reaction’s barrier height from RDB7 (Spiekermann, Pattanaik, and Green 2022b).

Split	MAE	RMSE
Random	3.87 ± 0.05	6.81 ± 0.28
Scaffold	6.28 ± 0.43	9.49 ± 0.50
K-means	5.47 ± 1.14	8.77 ± 1.85

Several approaches could be taken to further reduce the errors presented here. One could pre-train on additional data or fine-tune with experimental values. Ensembling is another established method to improve model predictions.

Related Software and Code Availability

In the machine learning space, `astartes` functions as a drop-in replacement for the ubiquitous `train_test_split` from `scikit-learn` (Pedregosa et al. 2011). Transitioning existing code to use this new methodology is as simple as running `pip install astartes`, modifying an `import` statement at the top of the file, and then specifying an additional keyword parameter. `astartes` has been especially designed to allow for maximum interoperability with other packages, using few dependencies, supporting all platforms, and validated support for Python 3.7 through 3.11. Specific tutorials on this transition are provided in the online documentation for `astartes`, which is available on GitHub.

Here is an example workflow using `train_test_split` taken from the `scikit-learn` documentation (Pedregosa et al. 2011):

```
import numpy as np
from sklearn.model_selection import train_test_split
```

```
X, y = np.arange(10).reshape((5, 2)), range(5)
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.33, random_state=42)
```

To switch to using `astartes`, from `sklearn.model_selection import train_test_split` becomes `from astartes import train_test_split` and the call to split the data is nearly identical and simple in the extensions that it provides:

```
import numpy as np
from astartes import train_test_split
```

```
X, y = np.arange(10).reshape((5, 2)), range(5)
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.33, sampler="kmeans", random_state=42)
```

With this small change, an extrapolative sampler based on k-means clustering will be used.

Inside cheminformatics, **astartes** makes use of all molecular featurization options implemented in AIMS_{Sim} (Bhattacharjee, Burns, and Vlachos 2023), which includes those from virtually all popular descriptor generation tools used in the cheminformatics field.

The codebase itself has a clearly defined contribution guideline and thorough, easily accessible documentation. **astartes** uses GitHub actions for Constant Integration testing including unit tests, functional tests, and regression tests. To emphasize the reliability and reproducibility of **astartes**, the data splits used to generate Table 1 and Table 2 are included in the regression tests. Test coverage currently sits at >99%, and all proposed changes are subjected to a coverage check and merged only if they cover all existing and new lines added as well as satisfy the regression tests.

Acknowledgements

The authors thank all users who participated in beta testing and release candidate testing throughout the development of **astartes**.

This material is based upon work supported by the U.S. Department of Energy, Office of Science, Office of Advanced Scientific Computing Research, Department of Energy Computational Science Graduate Fellowship under Award Number DE-SC0023112.

Disclaimer

This report was prepared as an account of work sponsored by an agency of the United States Government. Neither the United States Government nor any agency thereof, nor any of their employees, makes any warranty, express or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States Government or any agency thereof. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States Government or any agency thereof.

References

- Bannigan, Pauric, Matteo Aldeghi, Zeqing Bao, Florian Häse, Alan Aspuru-Guzik, and Christine Allen. 2021. "Machine Learning Directed Drug Formulation Development." *Advanced Drug Delivery Reviews* 175: 113806.
- Bemis, Guy W., and Mark A. Murcko. 1996. "The Properties of Known Drugs. 1. Molecular Frameworks." *Journal of Medicinal Chemistry* 39 (15): 2887–93.
- Bhattacharjee, Himaghna, Jackson Burns, and Dionisios G. Vlachos. 2023. "AIMS_{Sim}: An Accessible Cheminformatics Platform for Similarity Operations on Chemicals Datasets." *Computer Physics Communications* 283: 108579. <https://doi.org/https://doi.org/10.1016/j.cpc.2022.108579>.
- Bilodeau, Camille, Andrei Kazakov, Sukrit Mukhopadhyay, Jillian Emerson, Tom Kalantar, Chris Muzny, and Klavs Jensen. 2023. "Machine Learning for Predicting the Viscosity of Binary Liquid Mixtures." *Chem. Eng. J.*, 142454.
- Durdy, Samantha, Michael W. Gaultois, Vladimir V. Gusev, Danushka Bollegala, and Matthew J. Rosseinsky. 2022. "Random Projections and Kernelised Leave One Cluster Out Cross Validation:

- Universal Baselines and Evaluation Tools for Supervised Machine Learning of Material Properties.” *Digital Discovery* 1: 763–78.
- Géron, Aurélien. 2019. *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems*. O’Reilly Media, Inc.
- Heinen, Stefan, Guido Falk von Rudorff, and O. Anatole von Lilienfeld. 2021. “Toward the Design of Chemical Reactions: Machine Learning Barriers of Competing Mechanisms in Reactant Space.” *J. Chem. Phys.* 155 (6): 064105.
- Huyen, Chip. 2022. *Designing Machine Learning Systems: An Iterative Process for Production-Ready Applications*. O’Reilly Media, Inc.
- Jha, Swarn, Matthew Yen, Yazmin Salinas, Evan Palmer, John Villafuerte, and Hong Liang. 2023. “Learning-Assisted Materials Development and Device Management in Batteries and Supercapacitors: Performance Comparison and Challenges.” *Journal of Materials Chemistry A* 11: 3904–36.
- Jorner, Kjell, Tore Brinck, Per-Ola Norrby, and David Buttar. 2021. “Machine Learning Meets Mechanistic Modelling for Accurate Prediction of Experimental Activation Energies.” *Chem. Sci.* 12 (3): 1163–75.
- Komp, Evan, Nida Janulaitis, and Stéphanie Valteau. 2022. “Progress Towards Machine Learning Reaction Rate Constants.” *Physical Chemistry Chemical Physics* 24: 2692–2705.
- Lakshmanan, Valliappa, Sara Robinson, and Michael Munn. 2020. *Machine Learning Design Patterns: Solutions to Common Challenges in Data Preparation, Model Building, and Mlops*. O’Reilly Media, Inc.
- Landrum, Greg, and others. n.d. “RDKit: Open-Source Cheminformatics.” <https://www.rdkit.org>.
- Meredig, Bryce, Erin Antono, Carena Church, Maxwell Hutchinson, Julia Ling, Sean Paradiso, Ben Blaiszik, et al. 2018. “Can Machine Learning Identify the Next High-Temperature Superconductor? Examining Extrapolation Performance for Materials Discovery.” *Molecular Systems Design & Engineering* 3 (5): 819–25.
- Morgan, Harry L. 1965. “The Generation of a Unique Machine Description for Chemical Structures—a Technique Developed at Chemical Abstracts Service.” *Journal of Chemical Documentation* 5 (2): 107–13.
- Pedregosa, F., G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, et al. 2011. “Scikit-Learn: Machine Learning in Python.” *Journal of Machine Learning Research* 12: 2825–30.
- Ramakrishnan, Raghunathan, Pavlo O. Dral, Matthias Rupp, and O. Anatole von Lilienfeld. 2014. “Quantum Chemistry Structures and Properties of 134 Kilo Molecules.” *Scientific Data* 1 (1): 1–7.
- Ramsundar, Bharath, Peter Eastman, Patrick Walters, and Vijay Pande. 2019. *Deep Learning for the Life Sciences: Applying Deep Learning to Genomics, Microscopy, Drug Discovery, and More*. O’Reilly Media, Inc.
- Rogers, David, and Mathew Hahn. 2010. “Extended-Connectivity Fingerprints.” *Journal of Chemical Information and Modeling* 50 (5): 742–54.
- Ruddigkeit, Lars, Ruud Van Deursen, Lorenz C. Blum, and Jean-Louis Reymond. 2012. “Enumeration of 166 Billion Organic Small Molecules in the Chemical Universe Database GDB-17.” *Journal of Chemical Information and Modeling* 52 (11): 2864–75.
- Spiekermann, Kevin A., Lagnajit Pattanaik, and William H. Green. 2022a. “Fast Predictions of Reaction Barrier Heights: Toward Coupled-Cluster Accuracy.” *The Journal of Physical Chemistry A* 126 (25): 3976–86.

- . 2022b. “High Accuracy Barrier Heights, Enthalpies, and Rate Coefficients for Chemical Reactions.” *Scientific Data* 9 (1): 1–12.
- . 2022c. “High Accuracy Barrier Heights, Enthalpies, and Rate Coefficients for Chemical Reactions.” Zenodo. <https://doi.org/10.5281/zenodo.6618262>.
- Spiekermann, Kevin A., Lagnajit Pattanaik, William H. Green, Kevin Yang, Kyle Swanson, Wengong Jin, Connor Coley, et al. n.d. https://github.com/kspieks/chemprop/tree/barrier_prediction.
- Stuyver, Thijs, and Connor W. Coley. 2022. “Quantum Chemistry-Augmented Neural Networks for Reactivity Prediction: Performance, Generalizability, and Explainability.” *The Journal of Chemical Physics* 156 (8): 084104.
- Terrones, Gianmarco G., Chenru Duan, Aditya Nandy, and Heather J. Kulik. 2023. “Low-Cost Machine Learning Prediction of Excited State Properties of Iridium-Centered Phosphors.” *Chemical Science* 14: 1419–33.
- Tricarico, Giovanni A., Johan Hofmans, Eelke B. Lenselink, Miriam López Ramos, Marie-Pierre Dréanic, and Pieter FW Stouten. 2022. “Construction of Balanced, Chemically Dissimilar Training, Validation and Test Sets for Machine Learning on Molecular Datasets.” *10.26434/Chemrxiv-2022-M8l33*.
- Wang, Anthony Yu-Tung, Ryan J. Murdock, Steven K. Kauwe, Anton O. Oliynyk, Aleksander Gurlo, Jakoah Brgoch, Kristin A. Persson, and Taylor D. Sparks. 2020. “Machine Learning for Materials Scientists: An Introductory Guide Toward Best Practices.” *Chemistry of Materials* 32 (12): 4954–65.
- Wei, Jing, Xuan Chu, Xiang-Yu Sun, Kun Xu, Hui-Xiong Deng, Jigen Chen, Zhongming Wei, and Ming Lei. 2019. “Machine Learning in Materials Science.” *InfoMat* 1 (3): 338–58.
- Yang, Kevin, Kyle Swanson, Wengong Jin, Connor Coley, Philipp Eiden, Hua Gao, Angel Guzman-Perez, et al. 2019. “Analyzing Learned Molecular Representations for Property Prediction.” *Journal of Chemical Information and Modeling* 59 (8): 3370–88.
- Yang, Xin, Yifei Wang, Ryan Byrne, Gisbert Schneider, and Shengyong Yang. 2019. “Concepts of Artificial Intelligence for Computer-Assisted Drug Discovery.” *Chemical Reviews* 119 (18): 10520–94.